PyTorch Tutorial: Part II

Yunfei Teng

yt1208@nyu.edu

Department of Electrical and Computer Engineering
New York University Tandon School of Engineering

March 8, 2018

Outline

In this tutorial, some advanced models will be introduced.
For the codes of this tutorial, please check
https://github.com/yunfei-teng/PyTorch-Tutorial
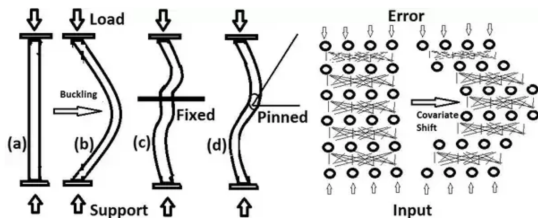
Normalization

Why do we need normalization?

- Batch Norm: Reduce internal co-variate shift.
- Instance Norm: Reduce the influence of contrast.

Actually, instance norm is a special case of batch norm($batchsize = 1$).

Batch Norm Paper: https://arxiv.org/pdf/1502.03167.pdf
Instance Norm Paper: https://arxiv.org/pdf/1607.08022.pdf

For both, Buckling or Co-Variate Shift a small perturbation leads to a large change in the later. *Debiprasad Ghosh, PhD, Uses AI in Mechanics*

Figure: Visual explanation from Quora: Co-Variate Shift

Normalization

Ideally, all the samples are i.i.d, but it's almost impossible. That's why we need normalization.
It not only accelerates the speed of convergence, but also increases the accuracy.

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1...m}\}$;
Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma,\beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^{m} x_i \qquad \text{// mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^{m} (x_i - \mu_{\mathcal{B}})^2 \qquad \text{// mini-batch variance}$$

$$\widehat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \qquad \text{// normalize}$$

$$y_i \leftarrow \gamma \widehat{x}_i + \beta \equiv \text{BN}_{\gamma,\beta}(x_i) \qquad \text{// scale and shift}$$

Figure: Core Algorithm of Batch Normalization

Autoencoder

Autoencoder is useful for information compression. Training an autoencoder to minimize reconstruction error amounts to maximizing a lower bound on the **mutual information** between input and learnt representation in latent space. [Stacked Denoising Autoencoders]


Figure: autoencoder architecture

Extended reading VAE: https://arxiv.org/pdf/1312.6114.pdf)

ResNet

ResNet Paper: `https://arxiv.org/pdf/1512.03385.pdf`

- Understand why the gradients are unstable. [gradient problem]
- Using residual blocks makes training arbitrarily deep neural nets become possible.

Figure: Residual block

Pre-trained Resnet models are provide in PyTorch:
`http://pytorch.org/docs/master/torchvision/`

Skip-connection

Effect of skip-connection: `https://arxiv.org/pdf/1712.09913.pdf`

(a) without skip connections

(b) with skip connections

Figure: The loss surfaces of ResNet-56 with/without skip connections.

Visual BackProp: Efficient visualization of CNNs

Paper: `https://arxiv.org/pdf/1611.05418.pdf`

Intuition: The feature maps contain less and less irrelevant information to the prediction decision when moving deeper into the network

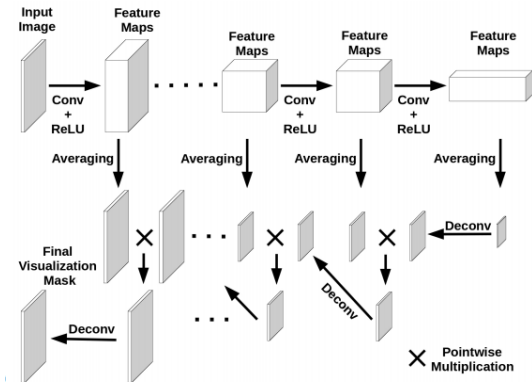Figure: VisualBackProp

Visual BackProp: Efficient visualization of CNNs

Figure: Visual Results

Generative Adversarial Nets

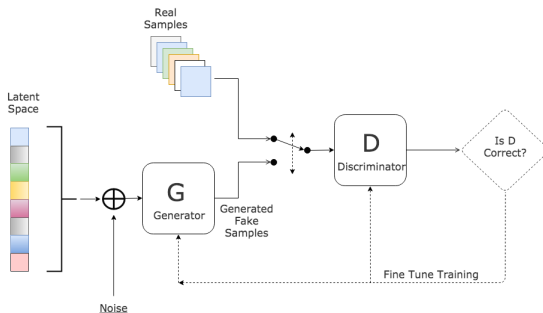**GAN paper**: `https://arxiv.org/pdf/1406.2661.pdf`

Figure: Generative adversarial nets

Extended Reading:
Nash Equilibria: `https://en.wikipedia.org/wiki/Nash_equilibrium`
Zero-sum game: `https://en.wikipedia.org/wiki/Zero-sum_game`